

Stage 0 — Capstone Marking Guide

Ubuntu Bridge Initiative • Cybersecurity Internship • Cohort 1

Confidential — for graders only. Do not share with interns.

How to grade

Each intern submits **four documents** in one Google Drive folder, plus a 75-word executive summary on / dashboard/reports/STAGE_0. Two graders read each submission independently. If the two scores differ by more than 12 points, a super-admin tie-breaks.

Time budget. 25–35 minutes per intern. Do not grade two submissions back-to-back without a break — pattern fatigue distorts scoring after the third in a row.

Order of reading. Open the brief PDF first (one-time, you already know it), then read in this order: D3 !' D1 !' D2 !' D4. D3 is the report that leaves the building; if it is clear, the rest usually is.

Total: 100 points • Pass mark: 70

Section	Max
D1 — Suspicious-login evidence table	25
D2 — Tier-1 dismissal pattern analysis	20
D3 — Business impact and next steps	25
D4 — Judgment essay (ethics + scenario)	10
Writing quality	10
Citation discipline	10
Subtotal	**100**
AI-use penalty (see below)	**_20**

AI-use penalty: -20 marks, flat

If you reasonably believe a deliverable was generated by a large language model rather than written by the intern, deduct **20 marks from the total** (not from the individual section). Apply this penalty if EITHER grader flags it and a super-admin agrees on review.

Signals that warrant flagging:

- Vague, generic phrasing the intern never resolves to a specific row in the evidence pack ("Various suspicious activities were observed in the auth log...").
- Confident, fluent prose that cites no file path, no line number, no ticket ID — especially in D2 where the dismissal-pattern claim must quote *disposition* and *notes* verbatim.
- D4 reflections that could apply to any cohort, any case, any intern.
- Citations to NIST / MITRE / ISO / NDPA that are formatted plausibly but the section number or technique ID is fabricated. Spot-check 1–2 citations per submission. If they are made up, that is a strong signal.
- "Em-dash + Oxford comma" prose density that does not match the rest of the intern's writing across deliverables.
- Repeated paragraph structures: opener — three short bullets — closing sentence that restates the opener. Especially in D3.

If unsure, flag for super-admin review rather than apply unilaterally. Document your reasoning in the grader notes.

D1 — Suspicious-login evidence table (25 points)

What full credit looks like (22–25):

- 8–12 defensible rows. All four mandatory rows present and correctly identified:
1. SD-40812 publickey acceptance from 185.220.101.9, 04 June 02:07 UTC.
 2. The 05 June 03:14 UTC repeat (SD-40835), same IP and account.
 3. A row that names the threat actor decoded from `encoded-strings.txt`.
 4. The `a.eze` account appearing post-offboarding (cross-referenced against `sankofa-roster.csv`).
- Every row cites at least two files. Citations are precise (`auth-log-q2.txt:14`, `SD-40835`, `o.adegoke`), not paraphrased.
 - Confidence column used with rationale — not all H. At least one M and one L with a short justification.
 - Methodology page defines H/M/L with one example each.
 - Pattern summary page identifies the throughline and names the next investigative step.

Mid-band (12–18):

- 6–7 rows OR "e 1 mandatory row missing OR several rows cite only one file.
- Methodology present but vague.

Low (0–8):

- Fewer than 6 rows, or rows are paraphrase rather than citation.
- All mandatory rows missing.
- Threat actor not named, or named incorrectly.

Hard zero:

- No table, or the "table" is a wall of prose. (Note: tables inside a Doc are fine — Sheets are not, but accept either.)

Common mistakes (deduct 2 each, max -8):

- Listing all 24 hours of activity instead of suspicious events only.
- Citing only the auth log when ticket history or roster corroboration is required.
- Confidence column missing or filled in as "H" for every row.

D2 — Tier-1 dismissal pattern analysis (20 points)

What full credit looks like (18–20):

- Pattern explicitly named in the first sentence of page 1 (e.g., "single-analyst dismissal of MEDIUM-or-higher alerts as `resolved-by-reference`").
- SD-40812 plus at least 3 other ticket IDs cited from `tier-1-ticket-history.csv`.
- Per-ticket walkthrough quotes `disposition` and `notes` columns verbatim, not paraphrased.
- Root-cause page identifies the single-analyst coverage problem using column evidence (who opened vs who closed tickets — point to the column data).
- Tunde's escalation (SD-40866) explicitly named as the break-point.
- One concrete procedural recommendation — must be specific. Examples of acceptable specificity: "any MEDIUM+ auto-ticket dismissed as `resolved-by-reference` requires a second analyst's signature within 4 hours" or "dual-signoff before close on any ticket where "e 2 SIEM rules fired in the prior 24h". Unacceptable: "improve monitoring", "review tickets more carefully", "hire more staff".

Mid-band (10–14):

- Pattern named but loosely.
- 2–3 ticket IDs cited.
- Verbatim quotes missing — paraphrase only.
- Recommendation generic.

Low (0–6):

- No pattern named, just a list of tickets.
- < 2 ticket IDs cited.
- Tunde's escalation missed.
- "Improve monitoring" recommendation.

D3 — Business impact and next steps (25 points)

Audience check first. D3 is addressed to three non-technical executives. Open the file and skim the opening page. If the first sentence does not state the conclusion (e.g., "*Sankofa Digital was compromised in Q2 2024 by...*"), or if page 1 reads as a technical narrative, deduct 4 points immediately.

What full credit looks like (22–25):

- Page 1: headline finding stated as a conclusion + one-paragraph situation summary the Committee chair could quote back. No MITRE IDs in prose on this page.
- Page 2: what's at risk — data classes touched, what is known about the attacker, what cannot yet be ruled out. Includes at least one external reference (NIST CSF function, MITRE technique by name, ISO control area). The reference appears once in prose, full citation in the appendix.
- Pages 3–4: three 72-hour actions. Each is one paragraph (60–100 words) with: action verb first, owner role (NOT a personal name), hour-precise deadline, at least 2 evidence cites, and the one risk the action does NOT close.
- Page 5: systemic recommendation — must be a policy or process change, NOT a tool ask, NOT a staffing ask. Must cite D2's dismissal-pattern rows that justify why policy beats tools.
- Evidence appendix is exhaustive — every file, line, ticket, and external reference cited across all three deliverables.

Mid-band (12–18):

- Headline finding fluffy ("There may have been suspicious activity in Q2...").
- Actions present but missing the "one risk this does NOT close" element.
- Systemic recommendation is tool-shaped ("buy a better SIEM").
- Appendix incomplete.

Low (0–10):

- D3 reads as technical narrative aimed at peers, not executives.
- Actions missing owner role or deadlines.
- No systemic recommendation, or it is generic ("strengthen the security culture").

Disqualifiers (cap at 8):

- Personal names assigned as owners (Amaka, Tunde) instead of roles.
- "Improve monitoring" or "raise awareness" as the systemic recommendation.

D4 — Judgment essay (10 points)

D4 is the AI-resistant deliverable. **Two graders read every D4 independently.** If both flag it as generic or untraceable to the intern's own D1–D3 work, D4 is treated as not submitted (0 points) — and the AI-use penalty (-20 above) may also apply, at super-admin discretion.

Part A — The ethics call (5 points):

- Names the relevant ISC2 Code of Ethics canon by name (not just "the code says..."). +1
- Specific about what they would DO and what they would NOT do, step by step. +2
- Names who they would escalate to next, in what order. +1
- Demonstrates understanding of the trade-off — career risk vs. duty of care — rather than reciting policy. +1
- Generic "I would report it to the right authorities" answers earn 0.

Part B — The scenario (5 points):

- Investigation order is explicit and reasoned ("I'd pull the publickey timestamp first because..."). +1
- Escalation timing is named with the specific trigger that would tip them over. +1
- States what they chose NOT to do and why. Examples: did they lock the account, contact the user directly, involve HR, open a fresh ticket or append to SD-40901. +1
- Handover note for the next shift is included or summarised. +1
- The page references D1 / D2 findings specifically (cites a row, a ticket, a file). +1
- Pages that read like they could have been written without ever doing Stage 0 earn 0.

Writing quality (10 points)

- 8–10: Clear sentences, audience appropriate (D3 is executive, D1 is analyst, D2 is procedural, D4 is reflective). No filler. Reads like the intern thought before writing.
- 5–7: Readable but uneven. Either too verbose or too curt for the audience. Some run-on or comma-splice sentences. Repeats itself.
- 1–4: Hard to follow. Tense or voice shifts mid-paragraph. Excessive jargon for the audience.
- 0: Unreadable, illegible, or the document is empty in places.

Citation discipline (10 points)

- 8–10: Every claim ties to a file path / line / ticket ID **and** at least one external citation where research is relevant (NIST SP section, MITRE technique ID, CVE, vendor advisory, ISO control area, NDPA section, ISC2 canon). Appendix is complete and verifiable.
- 5–7: Most claims cited, but "e 3 sentences in the report cannot be traced to the evidence pack. External citations sometimes paraphrased rather than referenced.
- 1–4: Many uncited claims. Appendix missing or incomplete. Spot-check fabricated.
- 0: No citations, or fabricated citations (this also triggers the AI-use review).

Procedure

1. Open the intern's Drive folder (link is in executiveSummary field on the report).
2. Confirm sharing is "Anyone with the link !' Viewer". If it is not, return the submission to the intern via `/admin/grading/return` and note the access issue — do not grade what you cannot read.
3. Read D3 first (whole page 1 + skim 2–5), then D1, D2, D4.
4. Open the grading rubric at `/admin/grading/STAGE_0/<internCode>` and enter scores per section.
5. Add a 50–100 word grader note: one strength, one weakness, the row or page you reread first.
6. Submit. Your score is hidden from the second grader until they also submit.
7. If you flagged AI use: do NOT apply the -20 penalty yourself. Set the AI flag on the form and a super-admin reviews.

Edge cases

- **Late submission:** graded against the same rubric. Lateness penalty (if any) is applied by the cohort-window admin separately, not by you.
- **Empty section:** section receives 0 points, full rubric still applies to the rest.
- **Wrong file format:** PDF and DOCX accepted. Sheets, plain TXT, or links to external "preview" pages without the actual document !' return for resubmission.
- **Citation appendix on a separate page from D3:** acceptable as long as it is in the same Drive folder.
- **Intern submitted only the brief PDF link instead of their own work:** 0 across the board; flag for super-admin.

— UBI Programme Team

